

Explorer

< | <> Untitled query ● > + Gemini settings ▾

Databases 1 Valid

ecommerce_db (Default)

Tables 2

▶ orders

▶ payments

Views 0

Events 0

Functions 0

Procedures 0

▶ information_schema

▶ mysql

▶ performance_schema

▶ sys

Queries 0

Run Save Format Clear

1 show tables;
2

Results Execution time: 2.3 ms

Tables_in_ecommerce_db

orders
payments

Rows per page: 20 ▾ 1 – 2 of 2

Google Cloud Bhanu Project 1 Search (/) for resources, docs, products, and more Search 5 ? B

All instances > e-commerce

Explorer Databases 1 ecommerce_db (Default) Tables 2 Views 0 Events 0 Functions 0 Procedures 0 information_schema mysql performance_schema sys

Run Save Format Clear

```
1 select user_id, count(order_id) as total_orders
2 from orders
3 group By user_id;
4
```

Results

user_id	total_orders
101	2
102	2
103	1

Execution time: 1.1 ms Export

Queries 0 Preview Rows per page: 20 1 - 3 of 3 < < > >>

The screenshot shows the Google Cloud Firestore Studio interface. On the left, there's a sidebar with various navigation options: Firestore, Database, Security, Indexes, Import/Export, Disaster Recovery, Time-to-live (TTL), Insights (Usage, Query insights, Monitoring, Key Visualizer), Release Notes, and a feedback icon.

The main area displays a hierarchical document structure under the path `/ > activity_logs > OldHHYTVQWw2PnXES9M`. The top-level node is `(default)`, which contains the `activity_logs` collection. This collection has a single document with the ID `OldHHYTVQWw2PnXES9M`. The document details are as follows:

- event:** "view_product"
- product_id:** "P006"
- timestamp:** December 22, 2025 at 2:28:00.000 A...
- user_id:** 102

Below the document details, there are five document IDs listed: `5t6mMBBcksDIDONw0qZp`, `6egyNH0nKM1Buav6n0SG`, `Phv90YaSz4dU60zXA1Vg`, `uzurQc4aaysaMxQdBHgF`, and `ye6cmbNljX2uca19e3Xy`.

Google Cloud Bhanu Project 1 firestore Search

Firestore Database (default)

Panel view Query builder

Query scope Collection /activity_logs

Limit * 100

Selection ORDER BY Field user_id Order descending

Add to query

Results Analysis

Query results

Document ID	event	product_id	timestamp	user_id
Phv90YaSz4dU6OzXA1Vg	"place_order"	"P005"	December 22, 2025 at 2:30:00 AM UTC+5:30	103
uzurQc4aaysaMxQdBHgF	"remove_from_cart"	"P006"	December 22, 2025 at 2:29:00 AM UTC+5:30	102
5t6mMBBcksDlDONw0qZp	"check_out"	"P002"	December 22, 2025 at 2:00:20 PM UTC+5:30	102
0ldHHYTVQWw2PnXTES9M	"view_product"	"P006"	December 22, 2025 at 2:28:00 AM UTC+5:30	102
ye6cmbNljX2uca19e3Xy	"add_to_cart"	"P001"	December 22, 2025 at 2:00:20 PM UTC+5:30	101
6egyNH0nKM1Buav6n0SG	"logout"	"P001"	December 22, 2025 at 2:31:00 AM UTC+5:30	101

Usage Query insights Monitoring Key Visualizer Release Notes

Why SQL is chosen for transactions:

1. When a customer places an order or makes payment, the data must be correct.
2. We cannot afford mistakes like double payment or missing order.
3. SQL databases give ACID guarantees, meaning the transaction fully completes or it rolls back.
4. Order and payment tables are related, so relational database fits well.
5. The structure of these tables does not change much, so fixed schema works.
6. SQL supports foreign keys, joins, constraints that protect integrity.
7. If two users order at the same time, SQL can handle concurrency safely.
8. Transaction history needs accuracy for audits, refunds, billing.
9. SQL writes are slower but reliable for financial operations.
10. That reliability makes SQL the preferred choice for money-related operations.

Why NoSQL is chosen for logs:

1. User activities generate huge amounts of data every second.
2. Logging every click to SQL would be slow and expensive.
3. NoSQL can handle high-speed continuous writes easily.
4. Logs do not need relationships or transactions.
5. Logs are mostly append-only, rarely updated.
6. Structure of logs can change anytime, NoSQL allows flexible schema.
7. Scaling NoSQL horizontally is easier and cheaper.
8. Storing large logs long-term costs less than SQL.
9. For analytics we mainly query latest N events, NoSQL supports indexed timestamps well.
10. NoSQL fits better for high volume, flexible, real-time logging workloads.