

Synopsis: Google Cloud Pub/Sub

- Google Cloud Pub/Sub is a fully managed messaging service used for real-time data communication.
- It follows the publish–subscribe (pub/sub) messaging model.
- Publishers send messages to a topic without knowing who will receive them.
- Subscribers receive messages asynchronously from the topic.
- It enables loose coupling between producers and consumers.
- Pub/Sub automatically scales to handle high message volumes.
- It provides high availability and fault tolerance.
- Messages are delivered with at-least-once delivery guarantee.
- Supports event-driven architectures and streaming data pipelines.
- Integrates easily with other Google Cloud services like Cloud Functions, Cloud Run, Big Query, and Dataflow.
- Useful for real-time analytics, log processing, notifications, and microservices communication.
- Reduces system dependency and improves performance and reliability.

TASK: Pull messages from a Pub/Sub topic and store them in a Big Query table using Cloud Functions Gen2.

1. Introduction

- Google Cloud Pub/Sub is a fully managed messaging service used for real-time event delivery.
- It enables asynchronous communication between distributed systems.
- This task uses Pub/Sub to stream events into Big Query for analysis.

2. Problem Statement

- Traditional systems struggle to handle real-time event ingestion at scale.
- Tight coupling between data producers and consumers reduces system flexibility.
- There is a need for a scalable, serverless solution to ingest and store event data.

3. Objective of the Task

- To capture event messages in real time using Pub/Sub.
- To process incoming messages using Cloud Functions Gen2.
- To store processed event data into Big Query for analytics.

4. System Architecture

- Publisher sends event messages to a Pub/Sub topic.
- Pub/Sub Topic acts as an event buffer.
- Cloud Function (Gen2) is triggered automatically on message arrival.
- Big Query stores the structured event data.

5. Working Principle

- Events are published to a Pub/Sub topic in JSON format.
- Pub/Sub delivers the message as a CloudEvent.
- Cloud Function decodes and processes the message.
- Extracted fields are inserted into a Big Query table with a timestamp.

6. Technologies Used

- Google Cloud Pub/Sub
- Cloud Functions (Gen2)
- Big Query
- Python
- Google Cloud IAM

7. Advantages of the System

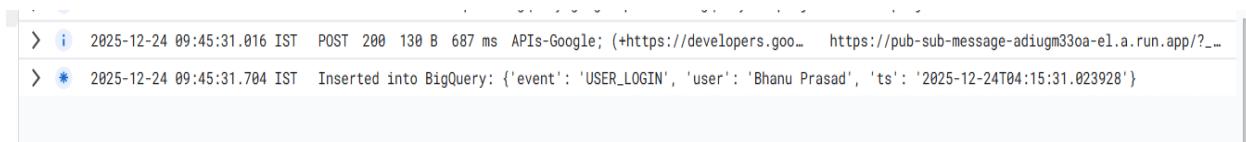
- Fully serverless and auto-scaling
- Loose coupling between components
- High reliability and fault tolerance
- Supports real-time analytics
- No infrastructure management required

8. Security and Permissions

- IAM roles control access to Big Query and Pub/Sub.
- Cloud Function uses a service account with limited permissions.
- Ensures secure and controlled data access.

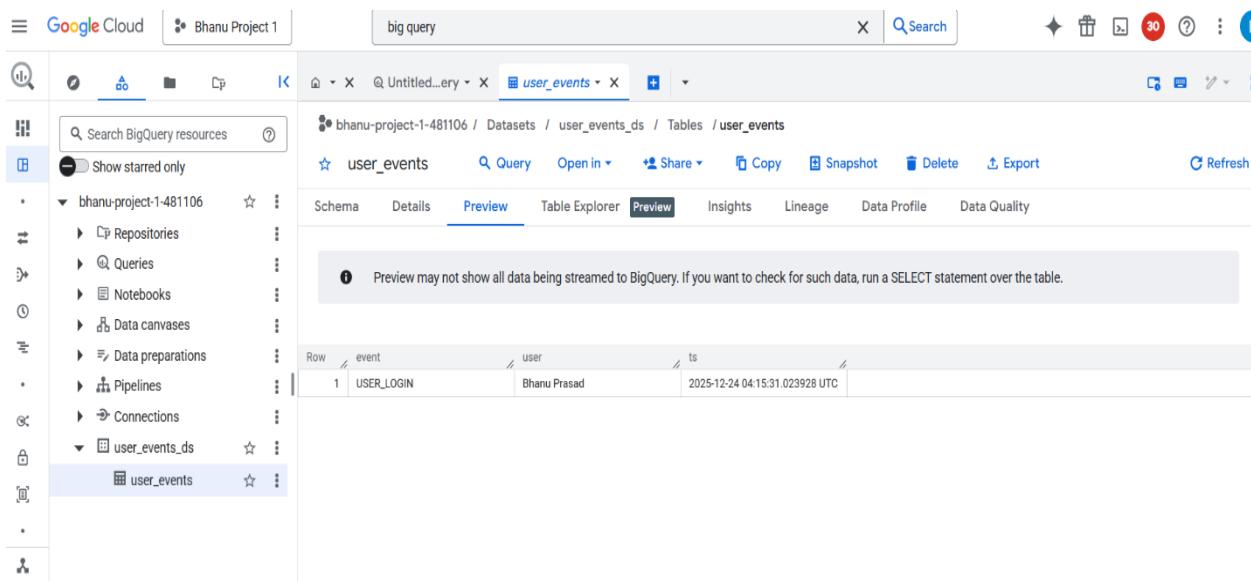
9. Implementation Results and Screenshots

- The Pub/Sub topic was successfully created and tested using sample event messages.
- Cloud Functions Gen2 was deployed and triggered automatically upon message publication.
- Event messages were decoded and processed without errors.
- Processed data was successfully inserted into the Big Query table.
- Screenshots are included to demonstrate:
 - Pub/Sub message publishing
 - Cloud Function execution logs
 - Successful data insertion into Big Query



A screenshot of a terminal window showing Cloud Function logs. It displays two log entries. The first entry is a POST request to a Pub/Sub endpoint. The second entry shows a successful insertion into a BigQuery table named 'user_events' with the data: {'event': 'USER_LOGIN', 'user': 'Bhanu Prasad', 'ts': '2025-12-24T04:15:31.023928'}

Cloud Function Logs



A screenshot of the Google Cloud Big Query interface. The left sidebar shows the project 'Bhanu Project 1'. The main area shows the 'user_events' table under the dataset 'user_events_ds'. The table has columns: Row, event, user, and ts. A single row is visible with the values: Row 1, event USER_LOGIN, user Bhanu Prasad, and ts 2025-12-24 04:15:31.023928 UTC.

Big Query Table

The screenshot shows the Google Cloud BigQuery interface. The left sidebar lists various projects and datasets, with 'mazenet-001' selected. In the main pane, a table named 'bhanu_user_events' is being previewed. The schema includes columns for 'event', 'user', and 'ts'. A single row is shown: 'Login_Successfully....' under 'event', 'Bhanu' under 'user', and '2025-12-24 04:33:12.024128 UTC' under 'ts'. The interface includes standard navigation and search tools.

The screenshot shows the Google Cloud Pub/Sub interface. On the left, the 'Topics' section is selected. A specific topic named 'user_events_topic' is displayed. The topic details include its name ('projects/bhanu-project-1-481106/topics/user_events_topic') and a list of subscriptions. One subscription is listed: 'user_events_topic-sub' with 'projects/bhanu-project-1-...' as the project. To the right, there is a detailed view of the topic's permissions, showing a list of principals with their roles. The permissions tab is active, showing roles like 'Editor' and 'Eventarc Service Agent'.

Pub/Sub Topic

10. Conclusion

- The Pub/Sub based architecture provides an efficient solution for real-time event ingestion.
- Integrating Pub/Sub with Cloud Functions and Big Query enables scalable data processing.